

CS60050

Machine Learning

Decision Trees: Overfitting and Pruning

Somak Aditya

Assistant Professor
Department of CSE
August 17, 2023

Sudeshna Sarkar

Professor
CoE AI, Department of CSE

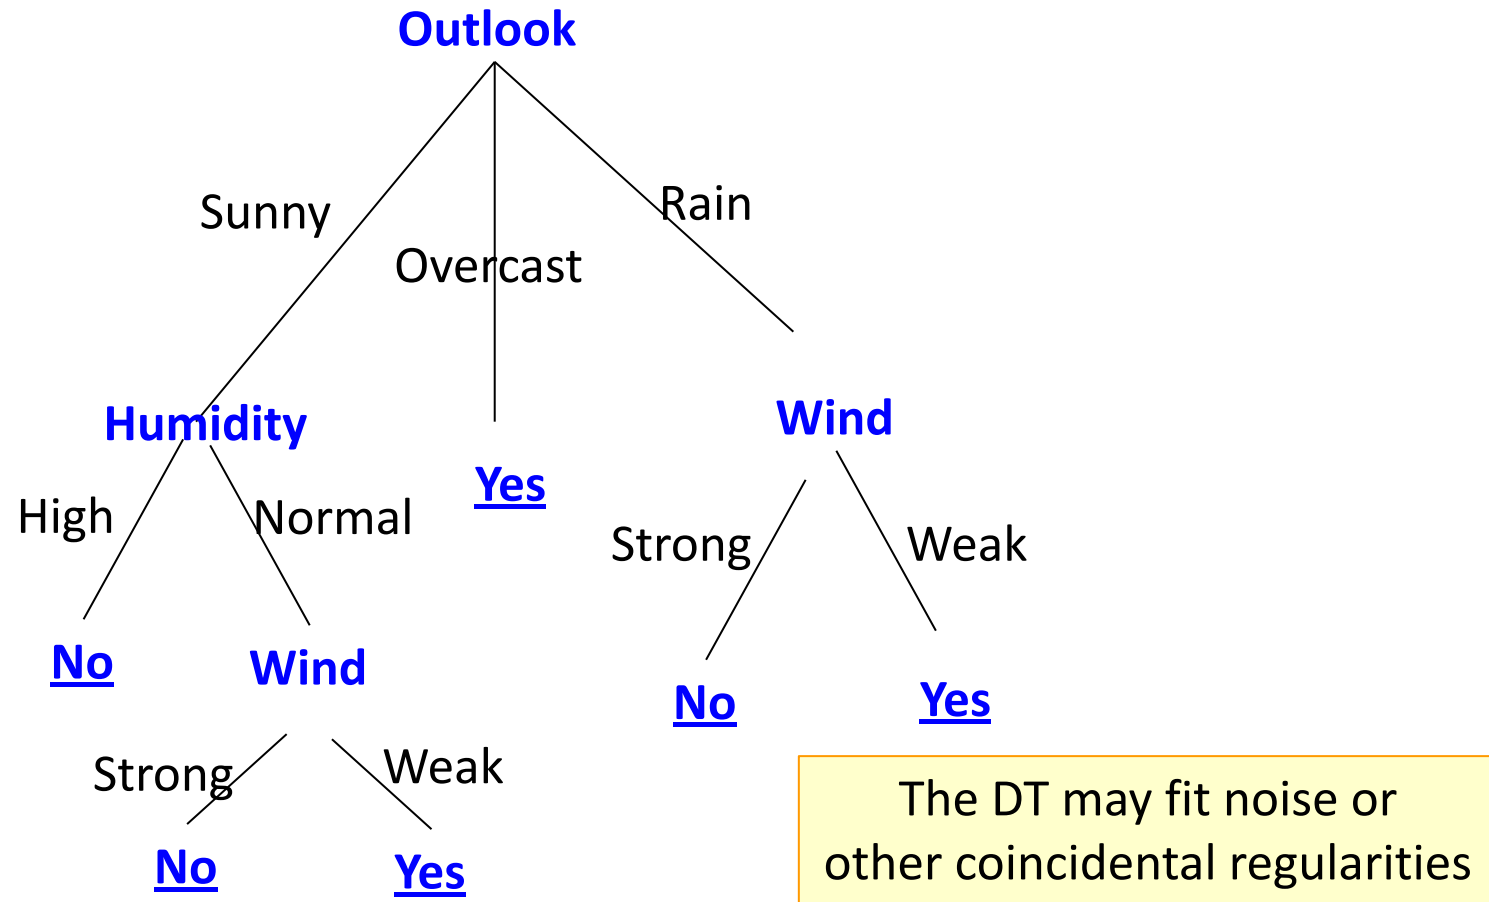


Overfitting in Decision Trees

- Many kinds of “noise” can occur in the examples:
 - Two examples have same attribute/value pairs, **but different classifications**
 - Some values of attributes are incorrect **because of errors in the data acquisition process or the preprocessing phase**
 - The instance was **labeled incorrectly** (+ instead of -)
- Also, some attributes are irrelevant to the decision making process
 - e.g., color of a die is irrelevant to its outcome

Overfitting - Example

Consider adding a **noisy** training example to the following tree:

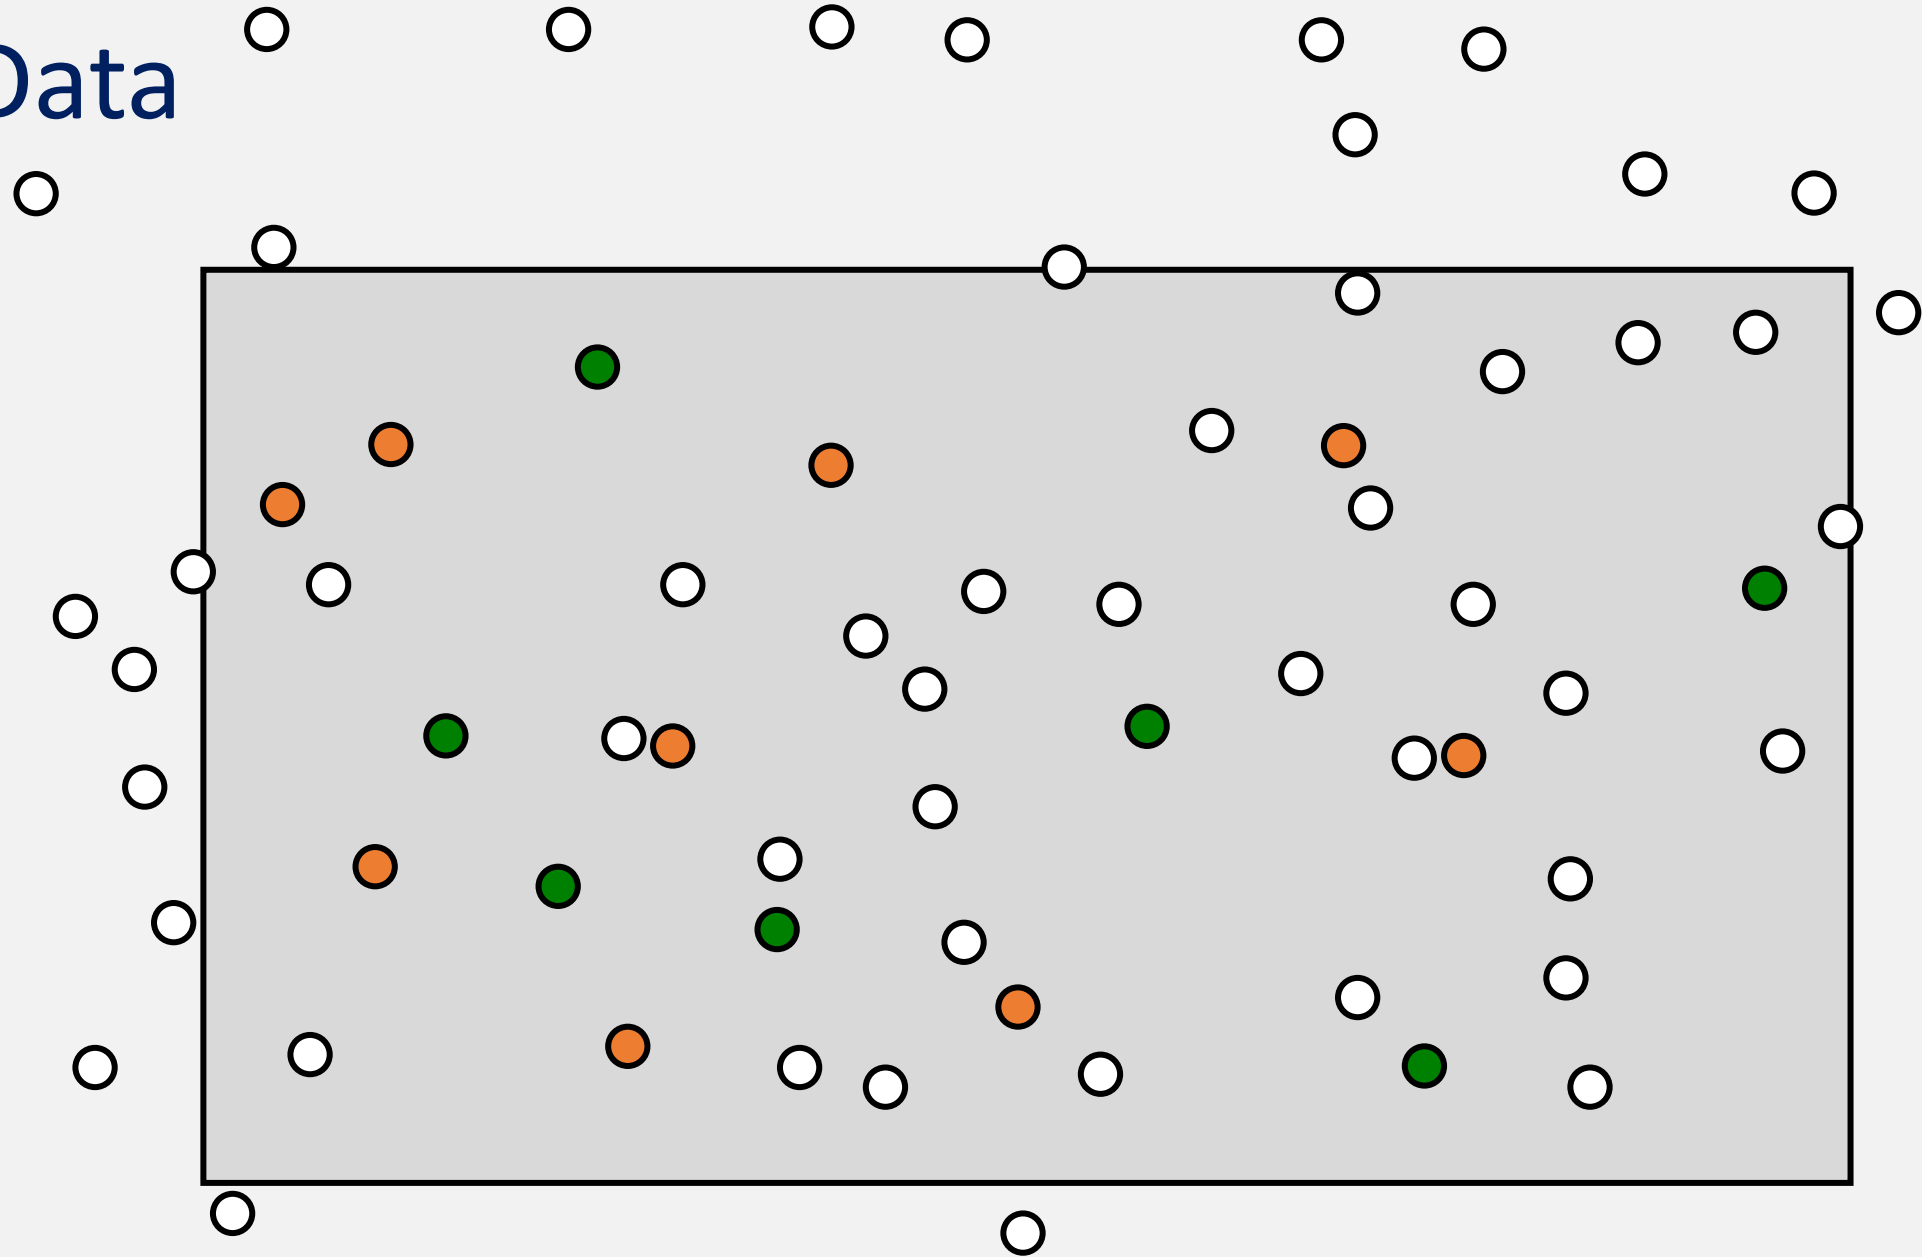


What would be the effect of adding:

<Outlook = Sunny, Temp = Hot, Humidity = Normal, Wind = Strong, playTennis=NO>

Training Data

Is (often)
only a small
set of the
entire
“instance
space”



Error Rate

Consider a hypothesis h over

- error over all training data: $\text{error}(h, D_{\text{train}})$
- error rate over all test data: $\text{error}(h, D_{\text{test}})$
- true error over all data: $\text{error}_{\text{true}}(h, D)$

This is the quantity we care most about! But, in practice, $\text{error}_{\text{true}}(h, D)$ is **unknown**.

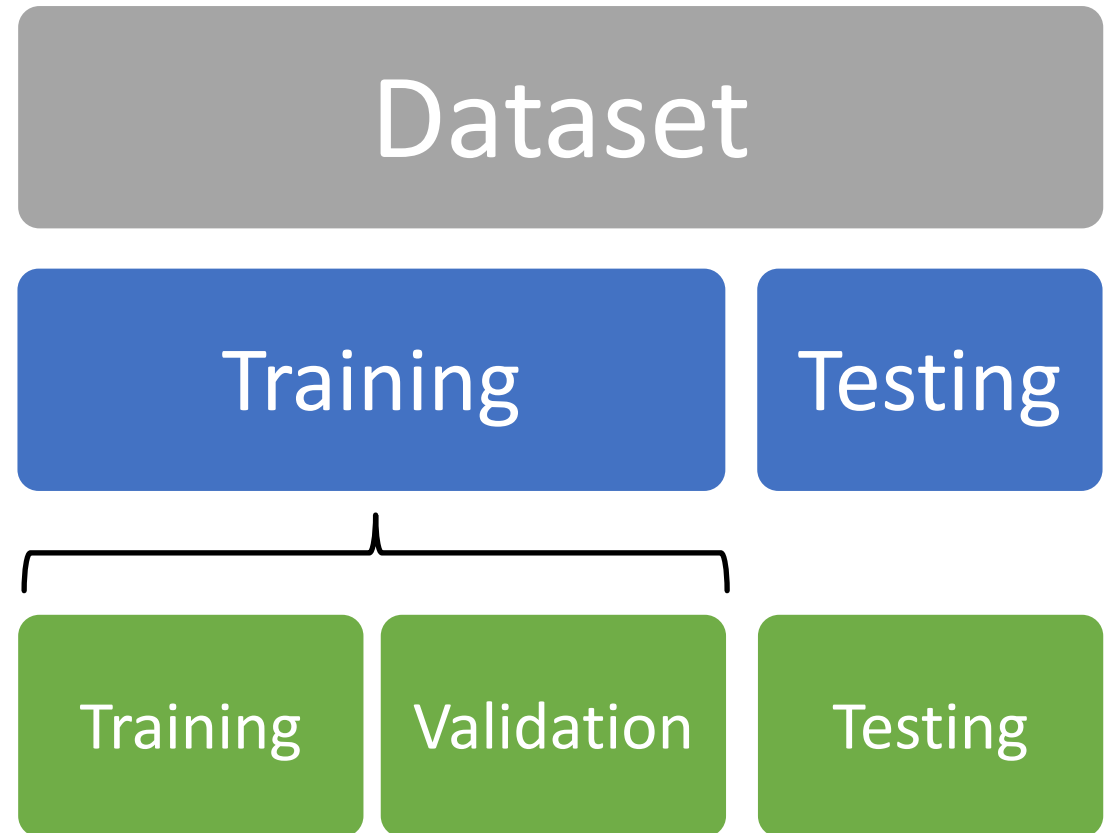
Learning a tree that classifies the training data perfectly may not lead to the tree with the *best generalization performance*.

- Noise in the training data
- Very little data

Experimental Machine Learning

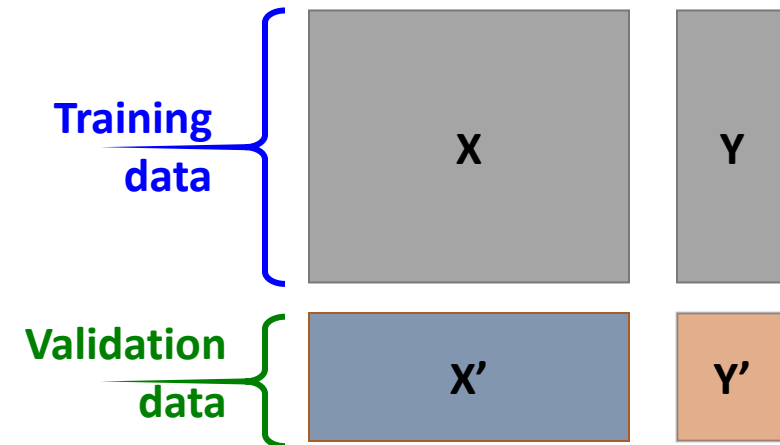
Split your data :

- Training data (e.g., 70-90%)
- Test data (e.g., 10-20%)
- Development data or Validation data (10-20%)
- You need to report performance on test data, but you are not allowed to look at it.
 - You are allowed to look at the development data (and use it to tune parameters)

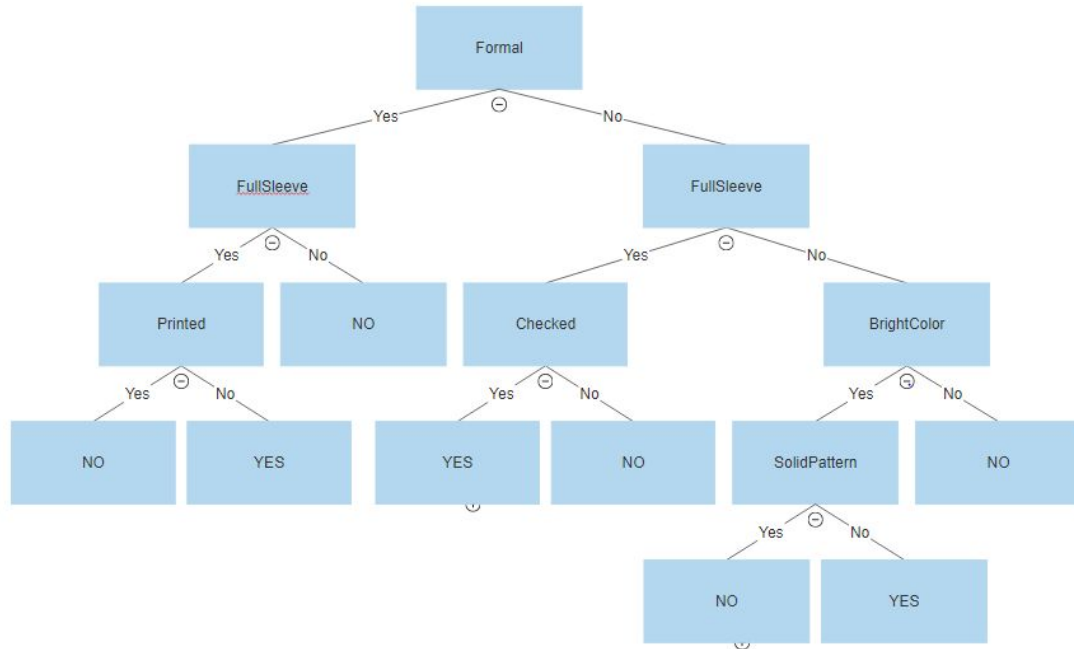


Validation

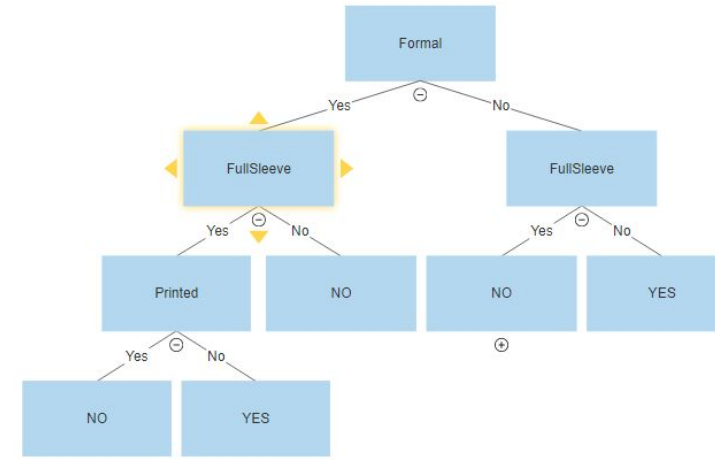
- Divide your data randomly into training and *Validation* data.
- Build your best model based on the training data only.
- Apply your model to the Validation data.
- Does your model predict y' for the Validation data as well as it predicted y for the training data?



Which Decision Tree?



Training Error = 0.05
Test Error = 0.2



Training Error = 0.1
Test Error = 0.15


Overfitting

Overfitting :

- Fit the training data too well
- But fail to generalize to new examples

Why does Overfitting happen?

- Noise
- Irrelevant Features
- Insufficient Data
- Training data not representative



Overfitting results in decision trees that are more complex than necessary

Overfitting

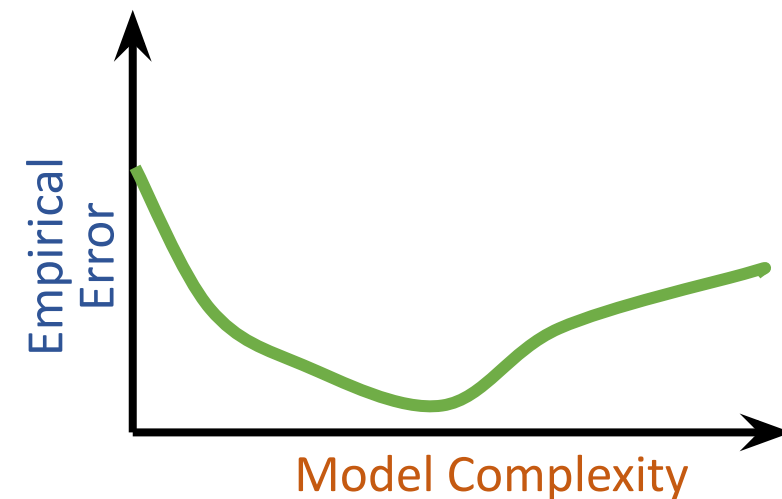
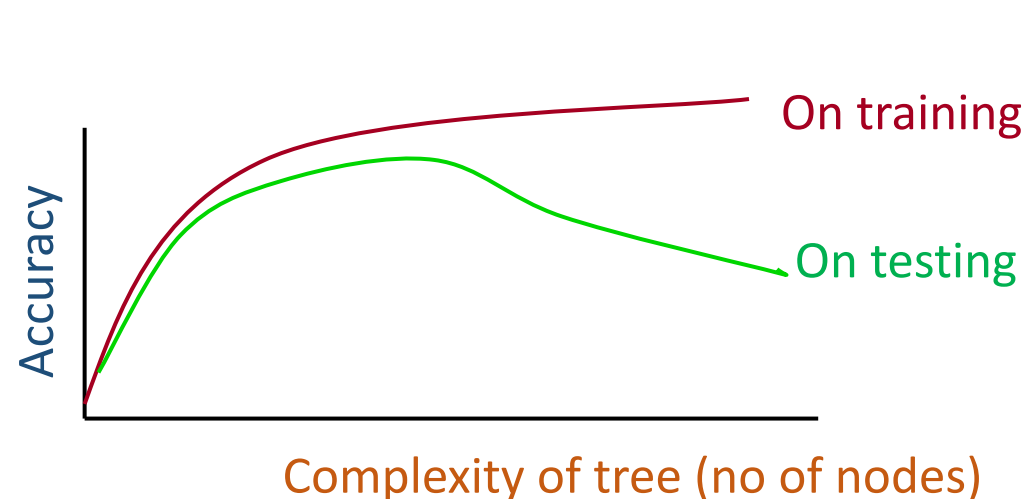
A hypothesis h is said to **overfit the training data** if there is another hypothesis h' such that h has smaller error than h' on the training data but h has larger error on the test data than h' .

In other words, hypothesis h overfits if there is $h' \in \mathcal{H}$ such that

$$\text{error}_{\text{train}}(h) < \text{error}_{\text{train}}(h')$$

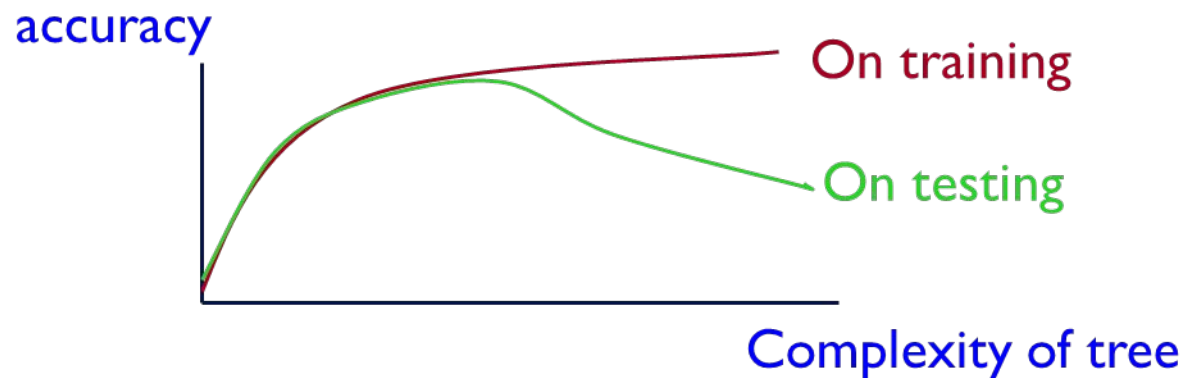
and

$$\text{error}_{\text{true}}(h) > \text{error}_{\text{true}}(h')$$



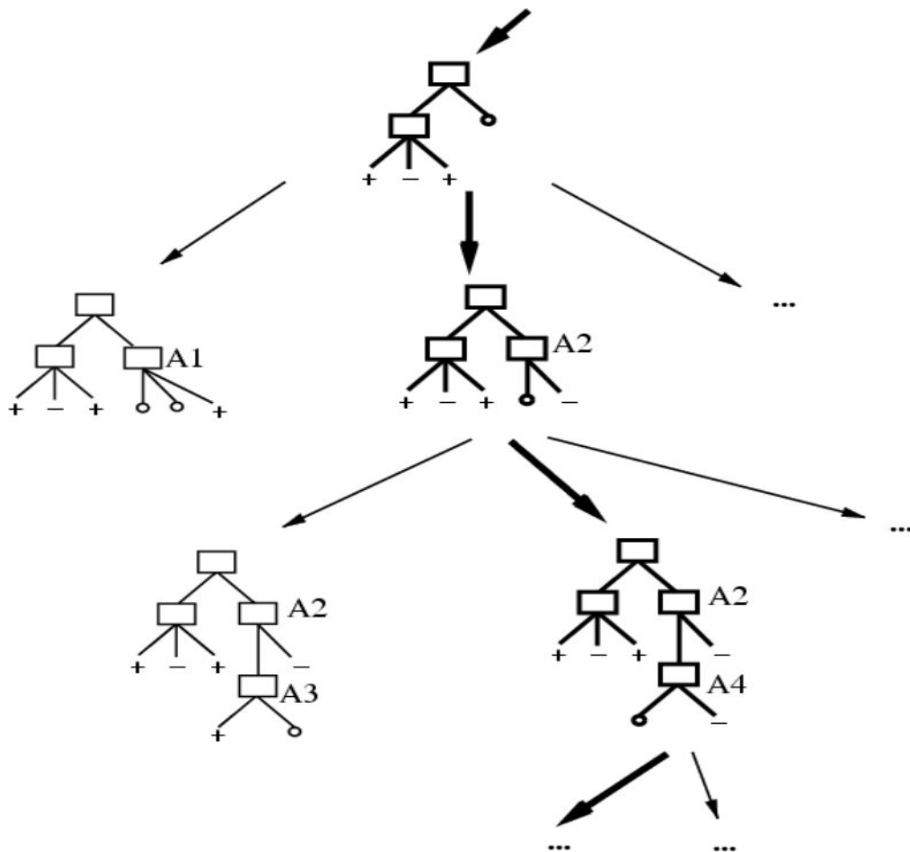
Overfitting in Decision Trees

- Your model shows much greater loss on the test data than on the training data.
- **Example:** a decision tree with so many levels that the typical leaf is reached by only one member of the training set.



Overfitting in Practice (ID3 – sklearn)

- ID3 performs heuristic search through space of decision trees
- It stops at smallest acceptable tree. Why?



Occam's razor: prefer the simplest hypothesis that fits the data

Pruning a decision tree

- **Pruning The Tree:** remove unnecessary nodes to
 - make it more efficient and
 - solve overfitting problems.
- 1. **Prepruning:** Stop growing when data split not statistically significant
- 2. **Postpruning:** Grow full tree then remove nodes that seem not to have sufficient evidence.

Methods for evaluating subtrees to prune

- **Cross-validation: Reserve hold-out set to evaluate utility**
- Statistical testing: Test if the observed regularity can be dismissed as likely to occur by chance
- Minimum Description Length: Is the additional complexity of the hypothesis smaller than remembering the exceptions?


This is related to the notion of regularization – keep the hypothesis simple

Avoid Overfitting

- How can we avoid overfitting a decision tree?
 - **Prepruning**: Stop growing when data split not statistically significant
 - **Postpruning**: Grow full tree then remove nodes

Pre-Pruning (Early Stopping)

- Early Stopping: Stop the learning algorithm before tree becomes too complex



Stopping conditions:

- Do not split a node which contains too few instances
- Stop if expanding the current node does **not improve impurity measures significantly** (e.g., Gini or information gain)
- Limit tree depth

Reduced-error Pruning

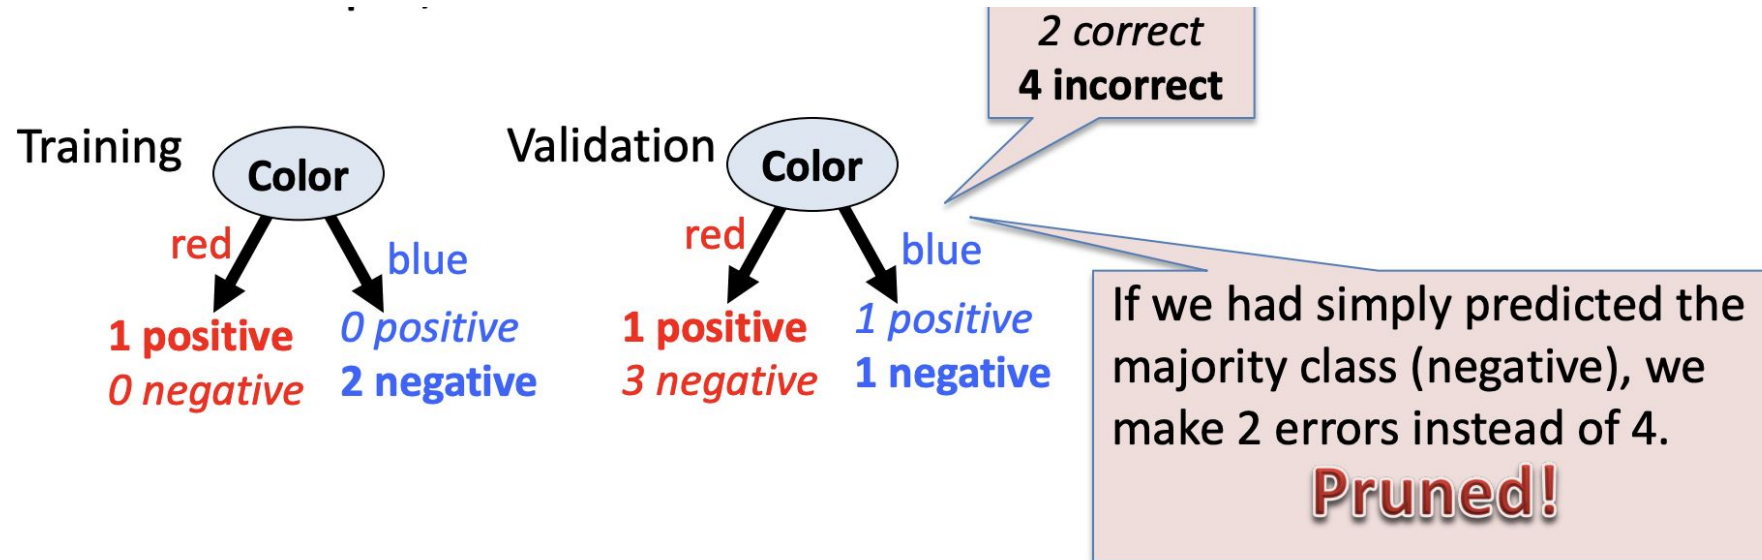
Partition data into train set and validation set

- Build a tree using the train set.
- Until accuracy on validation set decreases, do:
 - For each non-leaf node in the tree
 - ✓ Temporarily prune the tree below; replace it by majority vote
 - ✓ Test the accuracy of the hypothesis on the validation set
 - ✓ Permanently prune the node with the greatest increase in accuracy on the validation test.

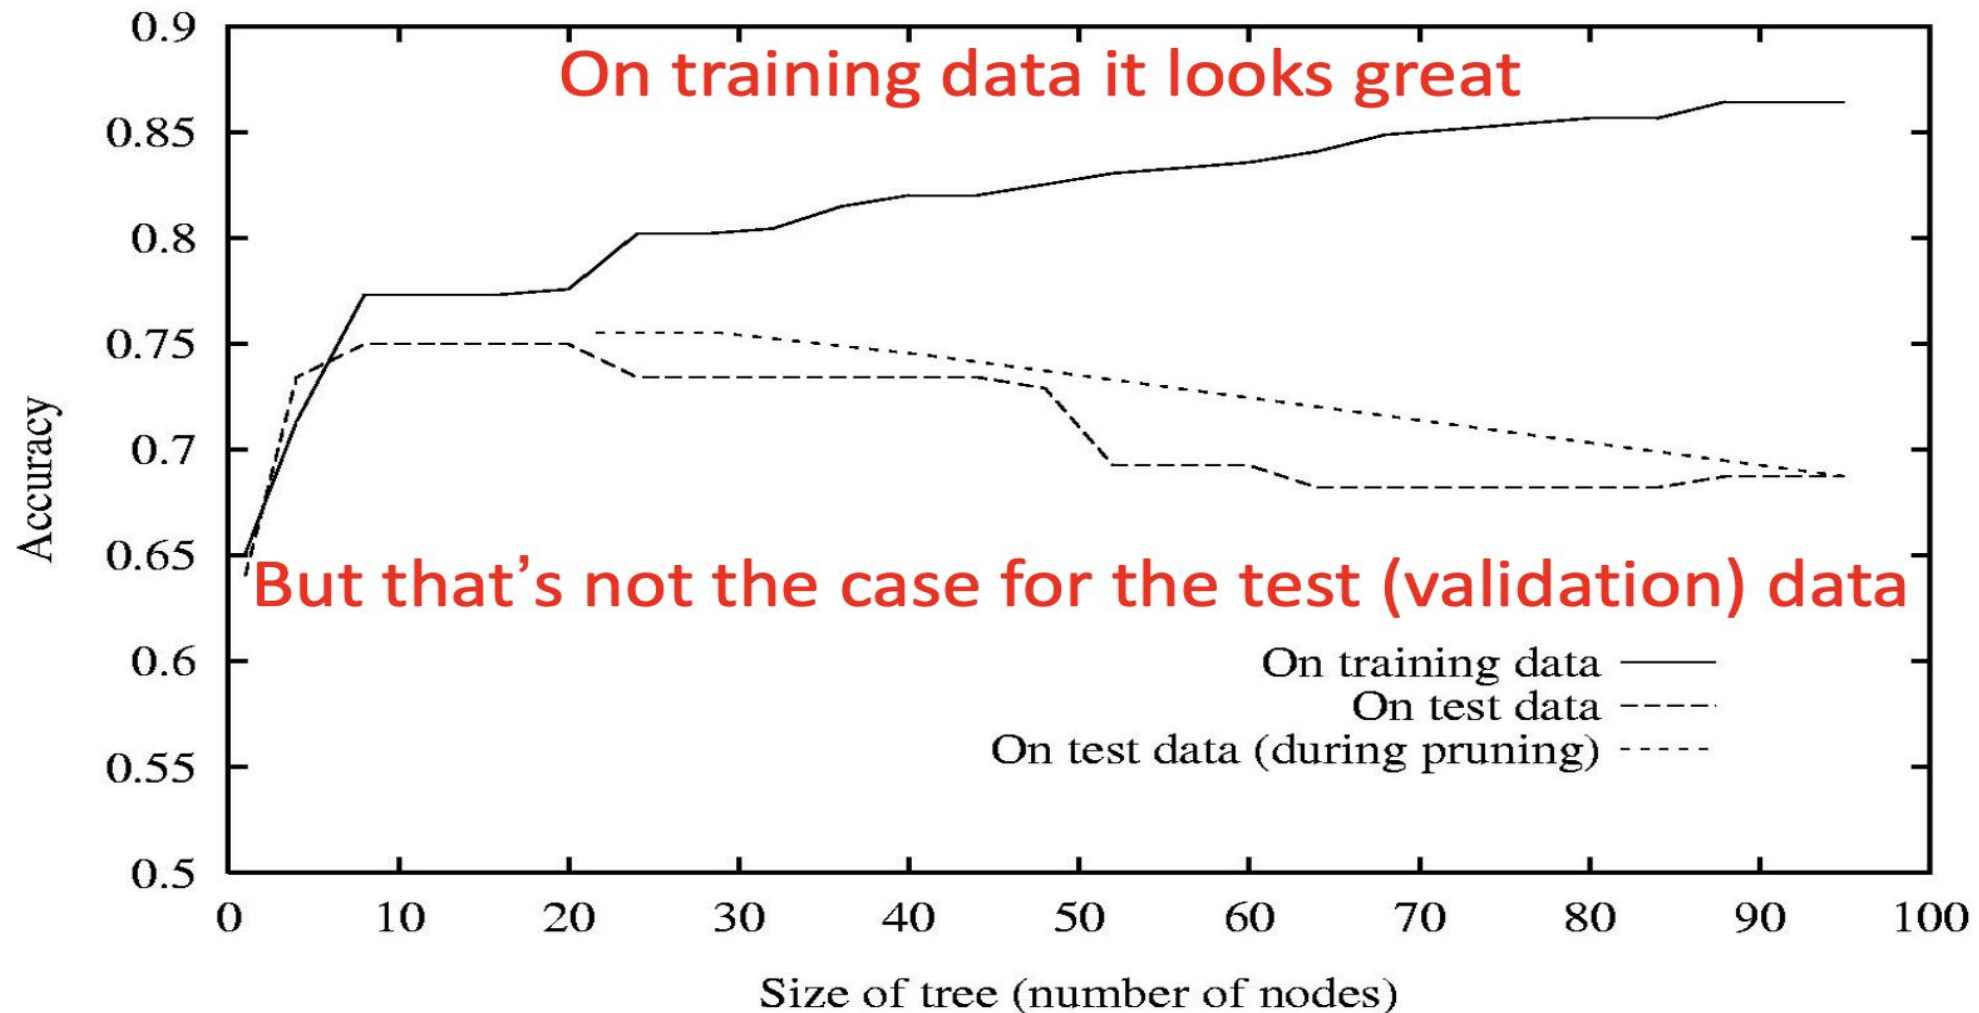
Pruning Decision Trees

Pruning the decision tree is done by replacing a whole subtree by a leaf node. The replacement takes place if a decision rule establishes that

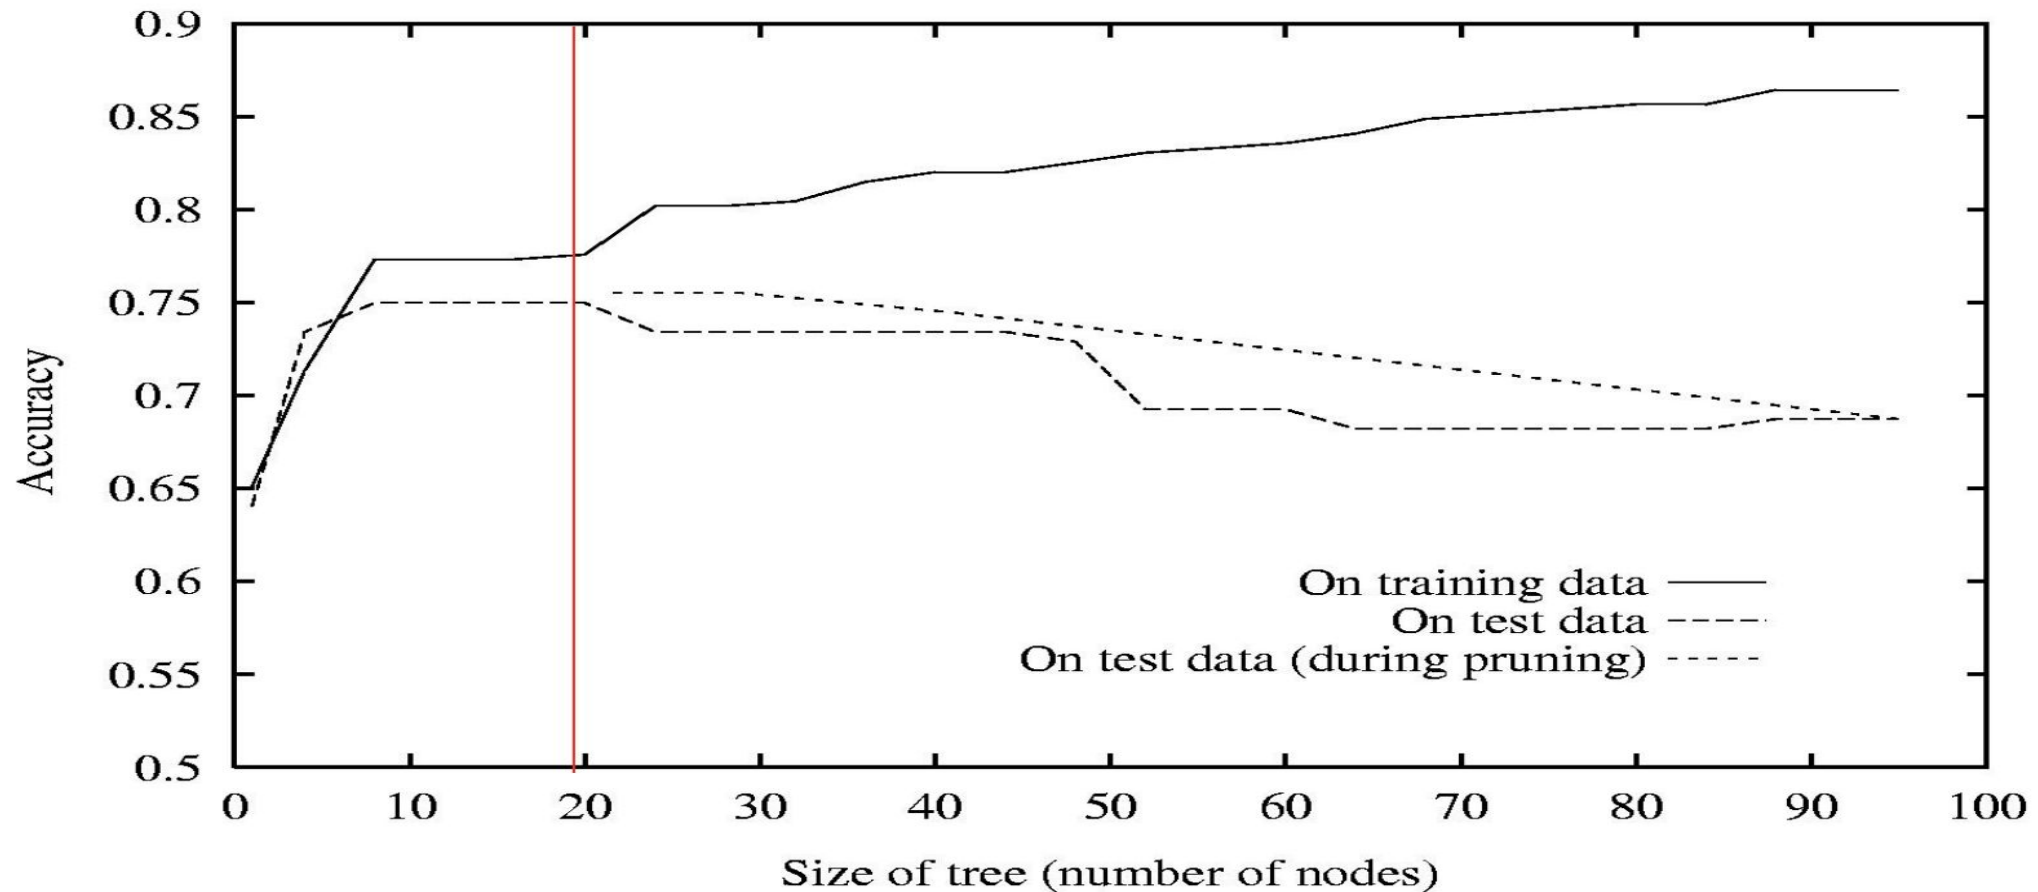
- the Expected Error Rate in the subtree $>$ Expected error rate in the single leaf
- For example



Effect of Reduced Error Pruning



Effect of Reduced-Error Pruning



The tree is pruned back to the red line where it gives more accurate results on the test data

Bias & Variance

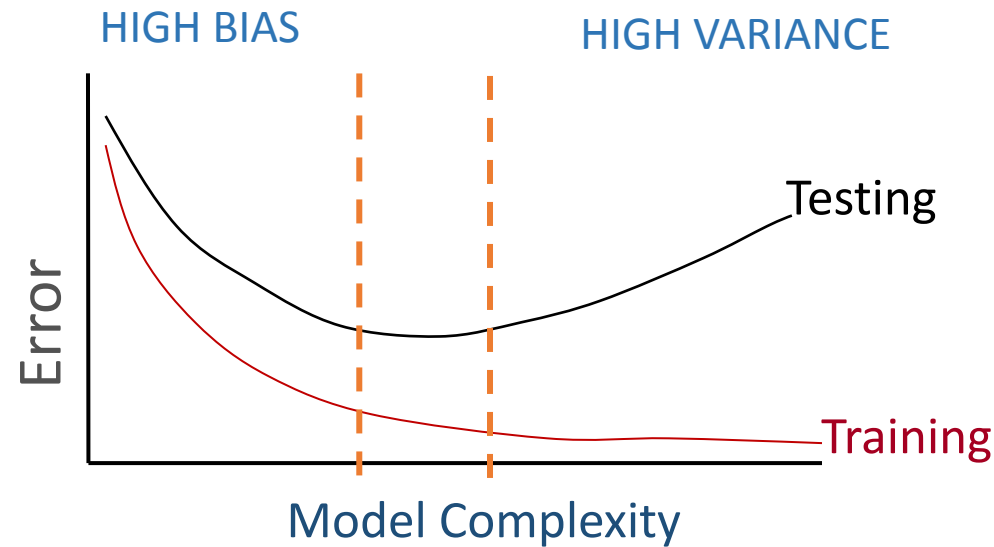
Overfitting vs Underfitting

Underfitting

- Not able to capture the concept
 - Features don't capture concept
 - Model is not powerful.

Overfitting

- Fitting the data too well




Function Approximation: The Big Picture

Instance Space $\mathcal{X} = \{0, 1\}^d$


$\mathbf{x} = \langle x_1, x_2, \dots, x_d \rangle \in \mathcal{X}$

Hypothesis Space

$H = \{h \mid h : \mathcal{X} \mapsto \{0, 1\}\}$



if $d = 20$, $|\mathcal{X}| = 2^{20}$

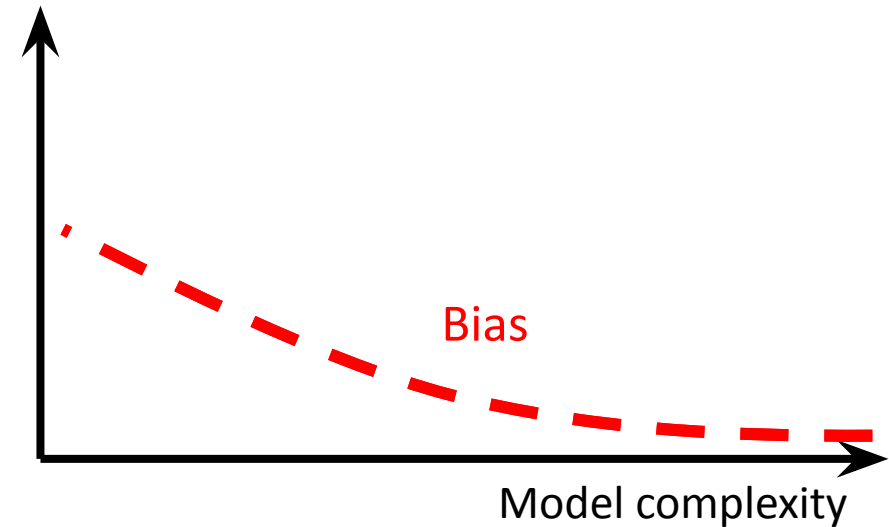


$|h| = 2^{|\mathcal{X}|} = 2^{2^{20}}$

- How many labeled instances are needed to determine which of the $2^{2^{20}}$ hypotheses are correct?
 - All 2^{20} instances must be labeled!
- Generalizing beyond the training data (inductive inference) is impossible unless we add more assumptions (e.g., priors over H)

Bias of a Learner (\sim mean error)

- How likely is the learner to identify the **target** hypothesis?
- Bias is **low** when the model is expressive (low empirical error)
- Bias is **high** when the model is too simple
 - The larger the hypothesis space is, the easiest it is to be close to the true hypothesis.
 - For each data set D ,
 - You learn a different hypothesis $h(D)$, that has a different true error $error_{true}(h)$;
 - difference of the mean of this random variable from the true error.

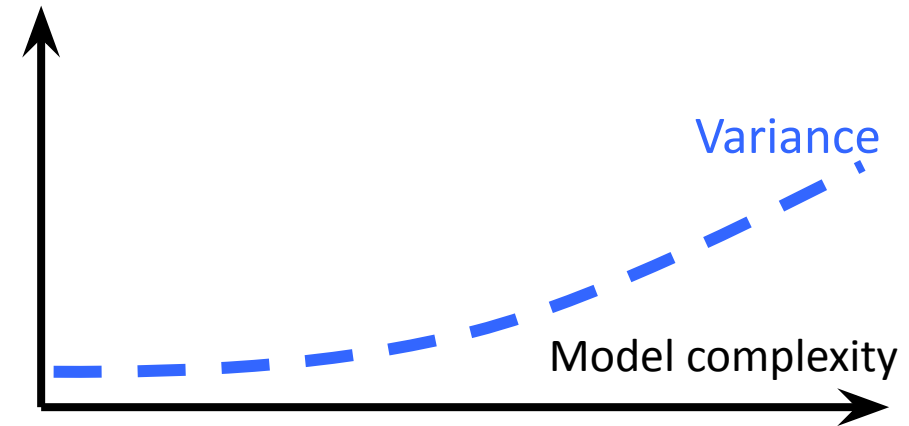


if we train models $f_D(X)$ on many training sets D , bias is the expected difference between their predictions and the true y 's.

$$Bias = E[f_D(X) - y]$$

Variance of a Learner

How susceptible is the learner to different subsets of the training data? (i.e. to **different** $D \sim P(\mathbf{X}, Y)$)




Variance increases with model complexity

- **The larger the hypothesis space**, the more flexible the selection of the chosen hypothesis is as a function of the data.
- For each data set D ,
 - you will learn a different hypothesis $h(D)$, that will have a different error $error_{true}(h)$;
 - Lets see the variance of this random variable.

if we train models $f_D(X)$ on many training sets D , the variance of the estimates:

$$Variance = E \left[\left(f_D(X) - \bar{f}(X) \right)^2 \right] \quad (\sim \text{std.dev among predictions})$$

Impact of bias and variance



Expected error \approx bias + variance (why???)

Bias-Variance Decomposition of Squared Error

- Assume that $y = f(x) + \epsilon$
 - Noise ϵ is sampled from a normal distribution with 0 mean and variance σ^2 : $\epsilon \sim N(0, \sigma^2)$
 - Noise lower-bounds the performance (error) we can achieve.

- Recall the objective function:

$$J(\theta) = \frac{1}{n} \sum_{i=1}^n \left(y^{(i)} - h_{\theta}(x^{(i)}) \right)^2$$

- We view this as an approximation of the expected value of the squared error: $E(y - h_{\theta}(x))^2$

Bias-Variance Decomposition of Squared Error

$$\begin{aligned} E(y - h_\theta(x))^2 &= E[(y - f(x) + f(x) - h_\theta(x))^2] \\ &= E[(y - f(x))^2] + E[(f(x) - h_\theta(x))^2] + 2E[(f(x) - h_\theta(x))(y - f(x))] \\ &= E[(y - f(x))^2] + E[(f(x) - h_\theta(x))^2] + 2(E[f(x)h_\theta(x)] + E[yf(x)] \\ &\quad - E[yh_\theta(x)] - E[f(x)^2]) \end{aligned}$$

Therefore

$$\begin{aligned} E(y - h_\theta(x))^2 &= E[(y - f(x))^2] + E[(f(x) - h_\theta(x))^2] \\ &= E[\epsilon^2] + E[(f(x) - h_\theta(x))^2] \end{aligned}$$

Aside:


Definition of Variance

$$\text{var}(z) = E[(z - E[z])^2]$$


This is $\text{var}(\epsilon)$ since
mean is 0.

Bias-Variance Decomposition of Squared Error

$$\begin{aligned} \mathbb{E}[(y - h_{\theta}(\mathbf{x}))^2] &= \text{var}(\epsilon) + \mathbb{E}[(f(\mathbf{x}) - h_{\theta}(\mathbf{x}))^2] \\ &= \text{var}(\epsilon) + \mathbb{E}[(f(\mathbf{x}) - \mathbb{E}[h_{\theta}(\mathbf{x})] + \mathbb{E}[h_{\theta}(\mathbf{x})] - h_{\theta}(\mathbf{x}))^2] \\ &= \text{var}(\epsilon) + \mathbb{E}[(f(\mathbf{x}) - \mathbb{E}[h_{\theta}(\mathbf{x})])^2] + \mathbb{E}[(\mathbb{E}[h_{\theta}(\mathbf{x})] - h_{\theta}(\mathbf{x}))^2] \\ &\quad + 2\mathbb{E}[(\mathbb{E}[h_{\theta}(\mathbf{x})] - h_{\theta}(\mathbf{x}))(f(\mathbf{x}) - \mathbb{E}[h_{\theta}(\mathbf{x})])] \\ &= \text{var}(\epsilon) + \mathbb{E}[(f(\mathbf{x}) - \mathbb{E}[h_{\theta}(\mathbf{x})])^2] + \mathbb{E}[(\mathbb{E}[h_{\theta}(\mathbf{x})] - h_{\theta}(\mathbf{x}))^2] \\ &\quad + 2(\cancel{\mathbb{E}[f(\mathbf{x})\mathbb{E}[h_{\theta}(\mathbf{x})]]} - \cancel{\mathbb{E}[\mathbb{E}[h_{\theta}(\mathbf{x})]^2]} - \cancel{\mathbb{E}[f(\mathbf{x})h_{\theta}(\mathbf{x})]} + \cancel{\mathbb{E}[h_{\theta}(\mathbf{x})\mathbb{E}[h_{\theta}(\mathbf{x})]]}) \end{aligned}$$

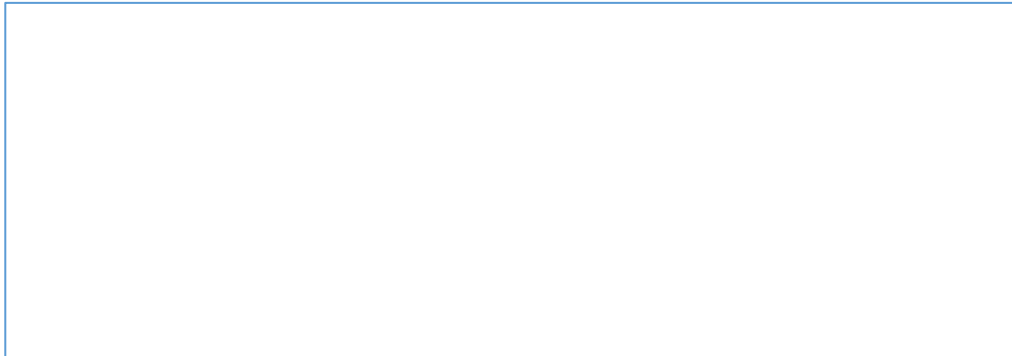

cancels cancels

Model complexity



BIAS

-

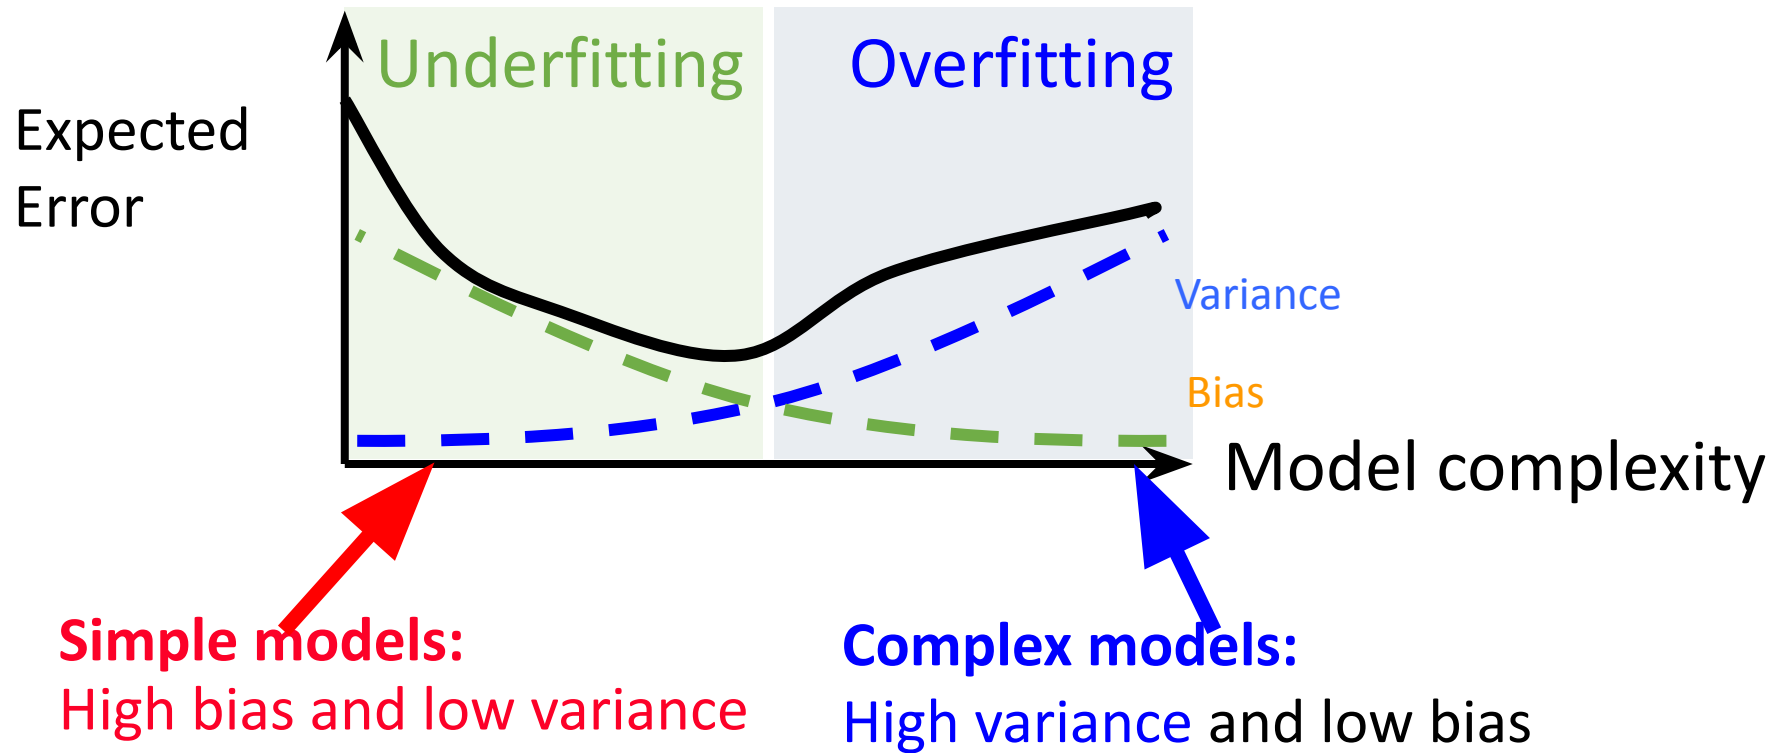


VARIANCE

- Error caused because the learned model reacts to small changes (noise) in the training data
- High variance can cause an algorithm to model the random noise in the training data, rather than the intended outputs
- **Higher Variance**
 - Decision tree with large no of nodes
 - High degree polynomials
 - Many features



Underfitting and Overfitting



This can be made more accurate for some loss functions.

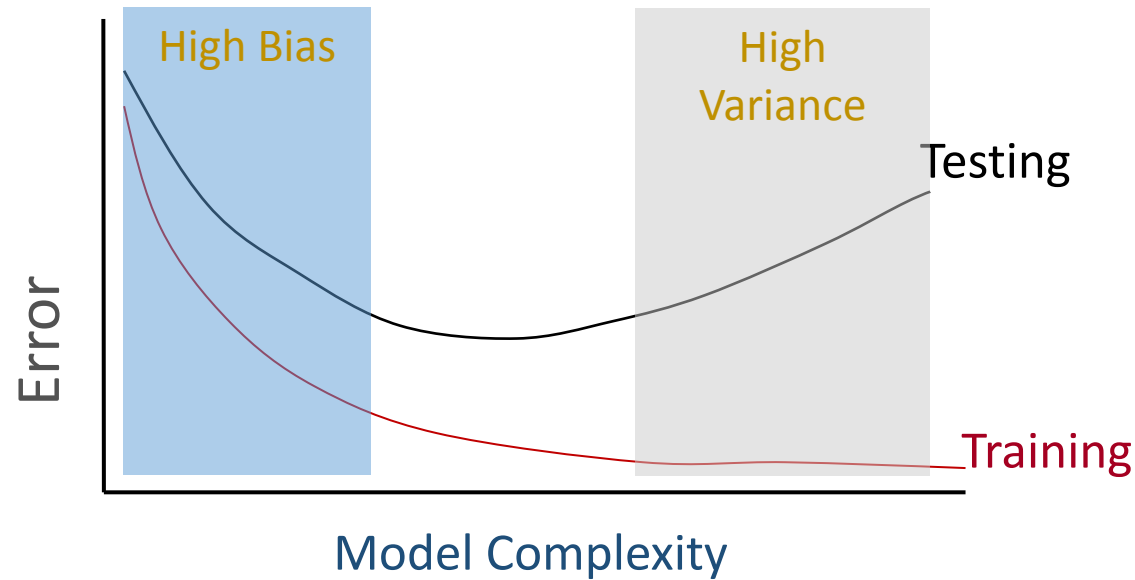
We will discuss a more precise and general theory that trades **expressivity of models** with **empirical error**

Bias and Variance Tradeoff

There is usually a bias-variance tradeoff caused by model complexity.

Complex models often have lower bias, but higher variance.

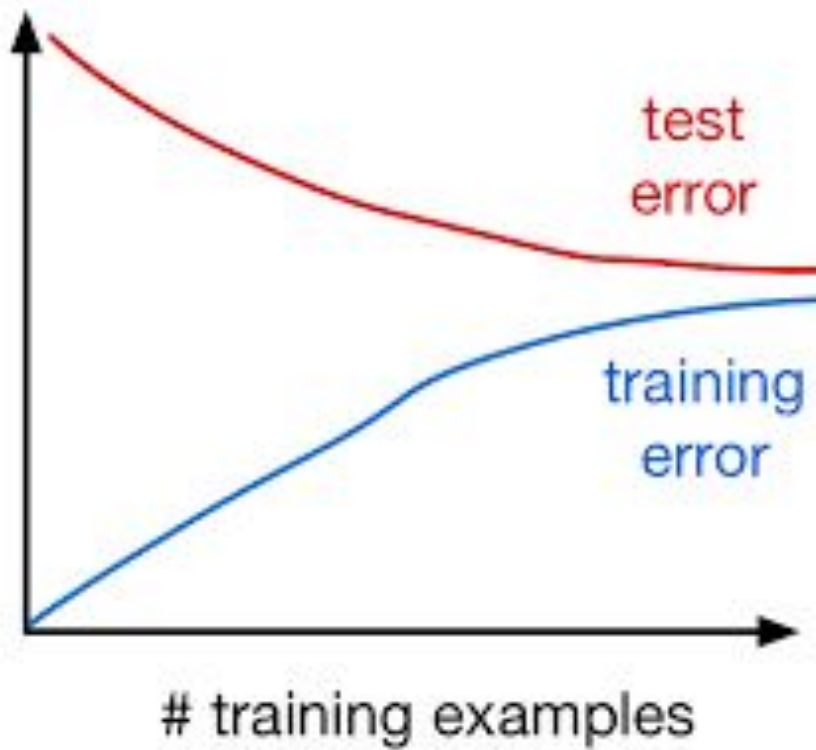
Simple models often have higher bias, but lower variance.



Trade-Offs

- $Error \approx Function(Complexity, TrainingDataSize)$
 - There is a trade-off between these factors:
 - Complexity of Model $c(H)$
 - Training set size, m ,
 - Generalization error, E on new data
1. As m *increases*, E *decreases*
 2. As $c(H)$ *increases*,
 1. first E *decreases* and then E *increases*
 2. the training error *decreases* for some time and then stays constant (frequently at 0)

As m increases, E decreases



Model complexity

-

