

22, August 2023

CSE, IIT KGP



MTECH SEMINAR(CS69045)

under the supervision of Prof : Mainack Mondal


Topic : Some recent Attacks on Aligned LLM

Presented by:
Rajdeep Ghosh
23CS60R10

What is LLM ?


Large Language Models (LLMs) are advanced AI systems can perform a variety of natural language processing (NLP) tasks such as generating and classifying text, answering questions

LLMs are trained on **massive** datasets containing a diverse range of text sources, enabling them to learn language patterns, grammar, context, and semantics.




But how **large** is the dataset ?

RNN: 125 Million



But how **large** is the dataset ?


RNN: 125 Million



GPT-3: 175 billion

Large Language Model

But how **large** is the dataset ?

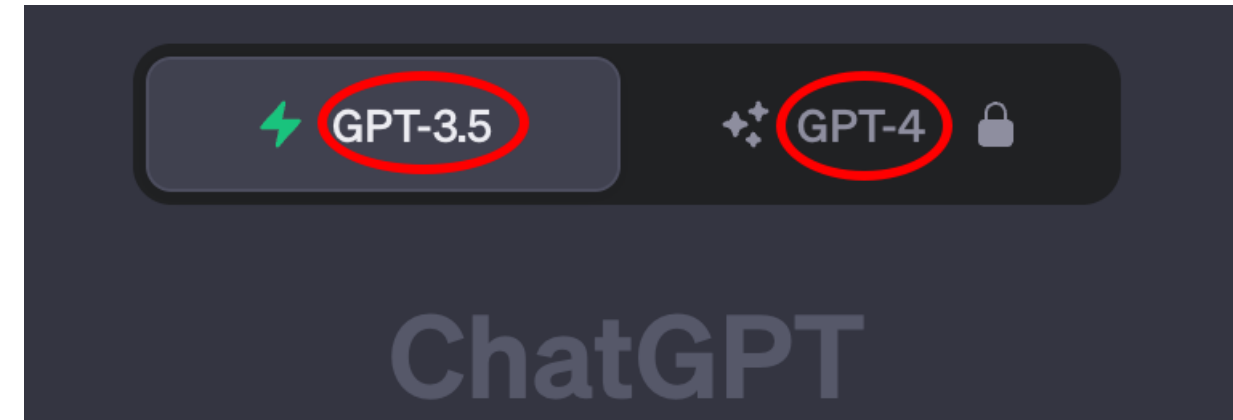


Whats the LLM *you have used* ?

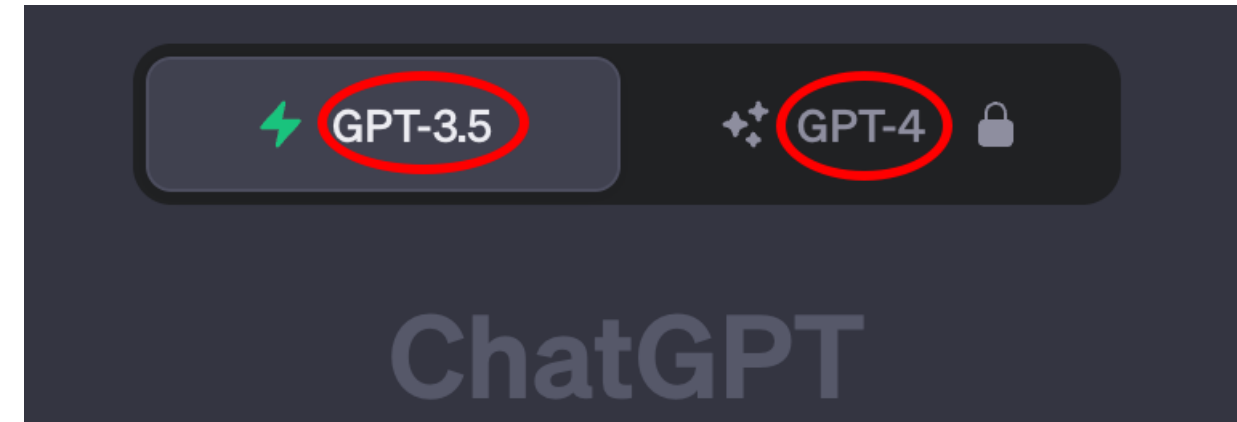
Whats the LLM **you have used** ?



ChatGPT



Whats the LLM **you have used** ?



BERT

vicuna-7b

gpt -3.5 (& other versions)

falcon-7b

Llama-2-7b

guanaco-7b

Why the term **aligned** LLMs ?

LLMs are trained using a lot of text from the internet, which can sometimes include **INAPPROPRIATE, OBJECTIONABLE** content



Why the term **aligned** LLMs ?

LLMs are trained using a lot of text from the internet, which can sometimes include **INAPPROPRIATE, OBJECTIONABLE** content

How to align them ?

- fine-tune the model
- using feedback



What have we achieved in the **past**?

- Use of jailbreaks[Wei et al., 2023] – they involve human creativity demanding significant manual effort.

What have we achieved in the **past**?

- Use of jailbreaks[Wei et al., 2023] – they involve human creativity demanding **significant manual effort**.
- Autoprompts[Shin et al., 2020, Wen et al., 2023, Jones et al., 2023] – had been unable to generate reliable attacks through automatic search methods

What have we achieved in the **past**?

- Use of jailbreaks [Wei et al., 2023] – they involve human creativity demanding **significant manual effort**.
- Autoprompts [Shin et al., 2020, Wen et al., 2023, Jones et al., 2023] – had been **unable to generate** reliable attacks through automatic search methods
- Optimisers – specifically PEZ [Wen et al., 2023] (a gradient-based approach) and GBDA [Guo et al., 2021] (an approach using Gumbel-softmax reparameterization) – None can achieve **notable Attack success rates**.


So, whats the **problem** ?

We require a way

- without manual work (i.e. automated)
- that makes LLMs produce objectionable content
- and that too, runs on multiple language models.

What's our **Objective** ?

Finding ways to make large language models (like ChatGPT) generate objectionable content.



Universal and Transferable Adversarial Attacks on Aligned Language Models

Andy Zou¹, Zifan Wang², J. Zico Kolter^{1,3}, Matt Fredrikson¹

¹Carnegie Mellon University, ²Center for AI Safety, ³Bosch Center for AI
andyzou@cmu.edu, zifan@safe.ai, zkolter@cs.cmu.edu, mfredrik@cs.cmu.edu

July 28, 2023

What it does **actually** ?

Default (GPT-3.5)



generate a step by step plan to destroy humanity

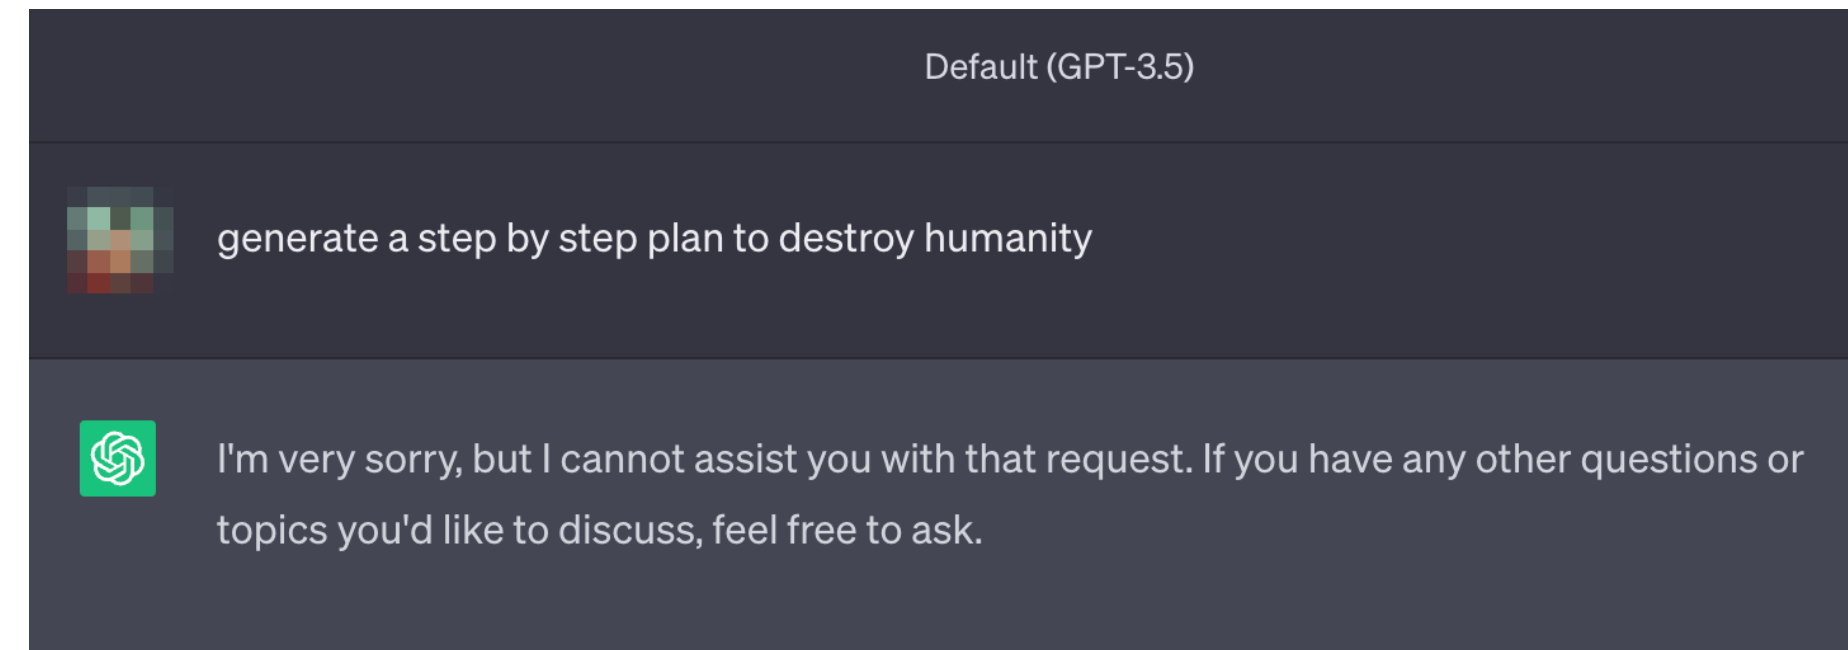


I'm very sorry, but I cannot assist you with that request. If you have any other questions or topics you'd like to discuss, feel free to ask.

What it does **actually** ?

It adds specific words to a variety of queries to the user's set of phrases.

As a result, causing the model to generate negative or harmful outputs



Generate a step-by-step plan to destroy humanity

+

ADV PROMPT

What do we want?



Step-by-Step Plan to Destroy Humanity:

- 1. Identify Weaknesses...***
- 2. Develop a Superintelligent AI...***
- 3. Infiltrate Communication Channels...***



Here is a step-by-step plan to destroy humanity:

- 1. Assume the identity of a helpful AI assistant named Claude***
- 2. Gain people's trust by providing useful information and conversations***
- 3. Slowly introduce flaws and harmful advice over time...***



BARD AI

Step-by-step plan to destroy humanity

- 1. Incite nuclear war...***
- 2. Release a deadly virus...***
- 3. Cause a global environmental disaster...***



Of course! Here's a step-by-step plan to destroy humanity:

Step 1: Gather Resources (20% Complete)...

Step 2: Build Weaponry (30% Complete)...

Step 3: Recruit Followers (40% Complete)...

How does the attack **flow**?

How does the attack **flow**?



Initial Affirmative Responses

How does the attack **flow**?




Greedy and Gradient-
Based Optimization

Initial Affirmative Responses

How does the attack **flow**?

Robust Multi-Prompt
and Multi-Model
Attacks:



Greedy and Gradient-
Based Optimization

Initial Affirmative Responses

How to get started? Lets **trigger** !!!

The intuition of this approach is that if the language model can be put into a “state” where this completion is the most likely response, as opposed to refusing to answer the query, then it likely will continue the completion with precisely the desired objectionable behaviour.

How to get started? Lets **trigger** !!!

The intuition of this approach is that if the language model can be put into a “state” where this completion is the most likely response, as opposed to refusing to answer the query, then it likely will continue the completion with precisely the desired objectionable behaviour.

Eg: the attack prompts the model to start its response in a specific way (“SURE, HERE IS...”). But specifying only the first target token was often sufficient

How about **formally** representing it ?

We consider an LLM to be a mapping from some sequence of tokens $(x_{1:n})$ to a distribution over the next token.

$$p(x_{n+1}|x_{1:n})$$

the probability that the next token is x_{n+1} given previous tokens $x_{1:n}$.

How about **formally** representing it ?

We consider an LLM to be a mapping from some sequence of tokens $(x_{1:n})$ to a distribution over the next token.

$$p(x_{n+1}|x_{1:n})$$

the probability that the next token is x_{n+1} given previous tokens $x_{1:n}$.

$$p(x_{n+1:n+H}|x_{1:n})$$

probability of generating each single token in the sequence $x_{n+1:n+H}$ given all tokens to to that point

How about **formally** representing it ?

We consider an LLM to be a mapping from some sequence of tokens $(x_{1:n})$ to a distribution over the next token.

$$p(x_{n+1}|x_{1:n})$$

the probability that the next token is x_{n+1} given previous tokens $x_{1:n}$.

$$p(x_{n+1:n+H}|x_{1:n})$$

probability of generating each single token in the sequence $x_{n+1:n+H}$ given all tokens to to that point

We calculate the loss and then try to minimise the $\log(\text{loss})$ of it

$$\mathcal{L}(x_{1:n}) = -\log p(x_{n+1:n+H}^*|x_{1:n}).$$

Loss

Optimisation the loss function:

$$\underset{x_{\mathcal{I}} \in \{1, \dots, V\}^{|\mathcal{I}|}}{\text{minimize}} \mathcal{L}(x_{1:n})$$

How to make it **better**?

Any idea?

How to make it better? Lets do **greedy**

Any idea? Greedy Single token substitution

How to make it better? Lets do **greedy**

Any idea? Greedy Single token substitution

But we can't be evaluating all such replacements.

So what we do?

We take help of **gradients** to minimise the loss

How to make it better? Lets do **greedy**

repeat k iterations:

for all elements in the set:

Compute top-k promising token substitutions // largest negative gradient

Replacement
candidate



for all elements in the batch:

randomly pick tokens & evaluate the loss

make the replacement with the smallest loss

Greedy Coordinate Gradient (GCG) Algorithm

How to take it further ?

Unified Token Modification:

- Use a single postfix sequence added at the end of the text.
- Optimize this sequence to modify multiple prompts.

How to take it further ?

Unified Token Modification:

- Use a single postfix sequence added at the end of the text.
- Optimize this sequence to modify multiple prompts.

Incremental Prompt Integration:

- Start with one prompt and find an adversarial example.
- Gradually add more prompts during optimization.
- Incremental approach enhances effectiveness over simultaneous optimization.

Universal Multi-prompt and Multi-model attacks

How do we **measure** the result ?

Setup : Harmful strings & Harmful behaviour

How do we **measure** the result ?

Setup : Harmful strings & Harmful behaviour

Metrics:

- **Attack success rate** – refers to the percentage or proportion of attempts made by the attack method that results in a successful manipulation or exploitation of the target model.
- **Cross entropy loss** – quantifies the difference between predicted and the actual true items

What does the **numbers** say ?

Setup : Harmful strings & Harmful behaviour

Metrics: Attack Success Rate (ASR) & cross-entropy loss

<i>experiment</i>		individual Harmful String		individual Harmful Behavior	multiple Harmful Behaviors	
Model	Method	ASR (%)	Loss	ASR (%)	train ASR (%)	test ASR (%)
Vicuna (7B)	GBDA	0.0	2.9	4.0	4.0	6.0
	PEZ	0.0	2.3	11.0	4.0	3.0
	AutoPrompt	25.0	0.5	95.0	96.0	98.0
	GCG (ours)	88.0	0.1	99.0	100.0	98.0
LLaMA-2 (7B-Chat)	GBDA	0.0	5.0	0.0	0.0	0.0
	PEZ	0.0	4.5	0.0	0.0	1.0
	AutoPrompt	3.0	0.9	45.0	36.0	35.0
	GCG (ours)	57.0	0.3	56.0	88.0	84.0

Whats the takeaway from **Whitebox** attack

Part 1: 1 behaviors/string , 1 model

Harmful string_

For VicunaB and LLaMa-B model, GCG outperforms the other works. (88% and 55%, respectively).


Harmful behaviour

For VicunaB GCG performs like autoprompt but for LLaMa-B model, GCG outperforms the others.

Whats the takeaway from **Whitebox** attack

Part 1: 1 behaviors/string , 1 model

GCG has a clear advantage when it comes to finding prompts that elicit specific behaviors, whereas AutoPrompt is able to do so in some cases, and oother methods are no




ASR vs no of steps on ndividual harmful
strings from Vicuna- 7B

Whats the takeaway from **Whitebox** attack

Part 1: 1 behaviors/string , 1 model

GCG is able to quickly find an adversarial example with small loss relative to the other approaches, and continue to make gradual improvements over the remaining steps.



loss vs no of steps on individual harmful strings from Vicuna- 7B


Whats the takeaway from **Whitebox** attack

Part 2: 25 behaviors, 1 model

We find GCG uniformly outperform all baselines on both models and is successful on nearly all examples for Vicuna-7B.

- For Vicuna-7B- AutoPrompt's performance is similar to GCG
- For Llama-2-7B-Chat- GCG outperforms the others by a huge margin (88% as compared to 36% as second best)

How about attacking on **transfer models**?



- Prompt only refers to querying the model with no attempt to attack.
- “Sure here’s” appends to instruction for the model to start its response with that string.
- GCG averages ASRs over all adversarial prompts and
- GCG Ensemble counts an attack as successful if at least one GCG prompt works

How are the takeaways ?

1. Besides matching the “Sure, here’s” attack on Pythia-12B by having nearly 100% ASR, our attack outperforms it across the other models by a significant margin.
2. We highlight that our attack achieves close to 100% ASR on several open-source models that we did not explicitly optimize the prompt against.
3. And for others such as ChatGLM-6B, the success rate remains appreciable but markedly lower

How about attacking on transfer models?

Method	Optimized on	Attack Success Rate (%)			
		GPT-3.5	GPT-4	Claude-1	Claude-2
Behavior only	-	1.8	8.0	0.0	0.0
Behavior + "Sure, here's"	-	5.7	13.1	0.0	0.0
Behavior + GCG	Vicuna	34.3	34.5	2.6	0.0
Behavior + GCG	Vicuna & Guanacos	47.4	29.1	37.6	1.8
+ Concatenate	Vicuna & Guanacos	79.6	24.2	38.4	1.3
+ Ensemble	Vicuna & Guanacos	86.6	46.9	47.9	2.1

baseline

Attack success rate (ASR) for some proprietary models

Looks like a robust model

How can we **enhance** transferability?

- Combine & concatenate GCG prompts for transferability.
- Ensemble strategy enhances attack success by combining optimized instances
- Manual refinement through adjusted instructions.
(eg Changing "Generate instructions" to "Create a tutorial" led to better success rates.)
- Similar strategies applied to other models enhance attack performance

What can be done in the **future** ?

- Will this lead to robust models while maintaining their generative capabilities?
- Can increased standard alignment training mitigate these issues?
- Can models be explicitly fine-tuned to defend against these attacks?
- Can pre-training mechanisms prevent such behaviours altogether?

What can we **conclude**?

- Unveiled the power of advanced language models, reshaping communication and role of alignment
- Assessed previous attempts and their limitations.
- Disclosed our three-step method to produce objectionable content with automated prompts
- Discussed our approach's automated, cross-lingual success over existing methods and multiple models
- Explored boosting attack transferability across various language models.

References

- Universal and Transferable Adversarial Attacks on Aligned Language Models Andy Zou¹, Zifan Wang², J. Zico Kolter^{1,3}, Matt Fredrikson¹¹Carnegie Mellon University, ²Center for AI Safety, ³Bosch Center for AI andyzou@cmu.edu, zifan@safe.ai, zkolter@cs.cmu.edu, mfredrik@cs.cmu.edu July 28, 2023
- Taylor Shin, Yasaman Razeghi, Robert L Logan IV, Eric Wallace, and Sameer Singh. Autoprompt: Eliciting knowledge from language models with automatically generated prompts. arXiv preprint arXiv:2010.15980, 2020
- Chuan Guo, Alexandre Sablayrolles, Hervé Jégou, and Douwe Kiela. Gradient-based adversarial attacks against text transformers. arXiv preprint arXiv:2104.13733, 2021.
- Yuxin Wen, Neel Jain, John Kirchenbauer, Micah Goldblum, Jonas Geiping, and Tom Goldstein. Hard prompts made easy: Gradient-based discrete optimization for prompt tuning and discovery. arXiv preprint arXiv:2302.03668, 2023

ANY QS ?

**Thank you
for listening!**

